

Service Level Indicators

Site Reliability Engineering



Objectives

In this module, we will look at the concept of service level indicator, and how it is different than a simple metric.

Learning Objectives

- Define service level indicator (SLI)
- Explain the importance of service level indicators
 - Business
 - Personnel



Service Level Indicator (SLI)


- ↘ Key to evaluating the performance of the business
- ↘ Quantitative measure of service
 - >>> Based on Service Level Objectives
- ↘ Business focused
 - >>> Service results that matter to the user
 - >>> Rises and falls along with user happiness
 - >>> Shows obvious difference during service outage
 - >>> Oscillates in a narrow band during normal operations

Common SLI Attributes

- ↘ Aggregates or percentiles relevant and meaningful
 - >>> SLO target
 - >>> Ranges match expectations
- ↘ Indicate progress in continuous improvement
- ↘ Measures related to
 - >>> User experience
 - >>> Complex system performance
- ↘ Need strong correlation between SLI and user happiness
 - >>> [Tune up your SLI metrics: CRE life lessons | Google Cloud Blog](#)

Service Level Indicator Focus

- ↘ Application availability and uptime
 - >>> Percent
- ↘ Request rates / throughput
 - >>> Compared to baseline
- ↘ Error rates
- ↘ Satisfaction
 - >>> Latency
 - >>> Availability
 - >>> Response time



Selecting SLI

- ↘ Identify system boundary
 - >>> Identify user-facing expectations at boundary
 - ~ What does availability look like
 - ~ How to measure
 - ~ Capture baseline
 - >>> Based on baseline and expectations
 - ~ Track as time-based data
 - ~ Confirm correlation of SLI with user expectations
- ↘ Fine tune SLI as business/user changes

SLI Specification Example

- ↘ User Transaction: Buying a Stock
- ↘ SLI Type: Availability
- ↘ SLI Specification: Proportion of buy page requests that are served successfully.
- ↘ SLI Implementations:
 - >>> Proportion of buy page requests that return a status code of 2xx.
 - ~ $\# \text{ buy page requests returning 2xx} / \text{total \# buy page requests} * 100$
 - ~ Should be greater than X
 - >>> Proportion of time that the OrderbookAPI container is accepting requests.
 - ~ $\text{successful page requests processed by orderbookapi} / \text{successful page requests through web server} * 100$
 - ~ Should be greater than Y

Actual Prometheus Metrics

Proportion of buy page requests that return a status code of 2xx

$$\left(\frac{\text{count}(\text{http_server_requests_seconds_bucket}\{\text{container}=\text{"orderbookapi"},\text{uri}=\text{"/buy"},\text{status}=\text{"200"}\})}{\text{count}(\text{http_server_requests_seconds_bucket}\{\text{container}=\text{"orderbookapi"},\text{uri}=\text{"/buy"}\})} \right) \times 1$$

Proportion of time that the OrderbookAPI container is accepting requests.

$$\left(\frac{\text{count}(\text{http_server_requests_seconds_bucket}\{\text{container}=\text{"orderbookapi"},\text{status}=\text{"200"}\})}{\text{sum}(\text{nginx_ingress_controller_requests}\{\text{exported_namespace}=\text{"orderbook-dev"},\text{status}=\text{"200"}\})} \right) \times 1$$

{mthree}

Example using Prometheus Metrics

↘ Success in using our sites ingress proxy

↘ Number of 200 successful requests in the last hour

```
count_over_time(nginx_ingress_controller_requests{exported_namespace="orderbook-dev",status="200"}[1h])
```

↘ Number of total requests in the last hour

```
sum(count_over_time(nginx_ingress_controller_requests{exported_namespace="orderbook-dev"}[1h]))
```

↘ Total successful requests compared to total requests

Example – Calculation

Ideally, we express this as a percentage

$$\left(\frac{("count_over_time(nginx_ingress_controller_requests\{exported_namespace = " @\"orderbook - dev\", status = \"200\"}[1h]) ")}{("sum(count_over_time(nginx_ingress_controller_requests\{exported_namespace = " @\"orderbook - dev\"}[1h])) ")} \right) \times 100$$

{mthree }

Frequently-Used KPIs

↘ Response Time (usually measured in minutes)

$$\frac{(When\ you\ respond\ to\ the\ alert) - (When\ the\ failure\ was\ detected)}{\#\ of\ Failures\ in\ a\ given\ time}$$

↘ Mean Time To Repair (MTTR)

$$\frac{Fixed\ DateTime - Time\ Detected}{\#\ of\ Failures\ in\ the\ given\ period}$$

↘ Average Cycle Time (measured for most important processes)

$$\frac{Finish\ Time - Start\ Time}{Total\ \#\ of\ Cycles}$$

SLI for People

- ↘ SLI not restricted to technology monitoring
- ↘ Can be used to justify SRE efforts and watch for negative trends
- ↘ Positive metrics
 - >>> Mean Time to Restore decreasing
 - >>> Mean Time to Detect decreasing
 - >>> Number of backlog issues decreasing
 - >>> Velocity of updates increasing
- ↘ Negative metrics
 - >>> Burn out flagged by increasing Recovery Time
 - >>> Runaway toil flagged by increasing issue backlog

Activity: SLI calculation in Prometheus/Grafana

Create the SLI metric in Prometheus or Grafana for the following:

- ↘ SLI Type:
 - >>> Latency
- ↘ SLI Specification:
 - >>> Orderbook API **buy** requests served within **300ms**
 - >>> Time period for this check is per **hour**
- ↘ Step-By-Step

Summary Q&A



{mthree}